

Assignment 4

Enhancing xv6 OS

Operating Systems and Networks

Monsoon 2021

Deadline: 26th October, 11:55 PM

There would be **NO** deadline extensions, so start early.

TOTAL MARKS: 100

Xv6 is a simplified operating system developed at MIT. Its main purpose is to explain the main concepts of the operating system by studying an example Kernel. xv6 is a re-implementation of Dennis Ritchie's and Ken Thompson's Unix Version 6 (v6). xv6 loosely follows the structure and style of v6, but is implemented for a modern RISC-V multiprocessor using ANSI C. **You will be tweaking the Xv6 operating system as a part of this assignment.** You can download the xv6 source code from [here](#). You can see the install instructions [here](#).

Specification 1: syscall tracing [15 Marks]

In this assignment, you will add a system call, trace, and an accompanying user program strace.

```
strace mask command [args]
```

strace runs the specified command until it exits.

- It intercepts and records the system calls which are called by a process during its execution.
- It should take one argument, an integer mask, whose bits specify which system calls to trace.

For example, to trace the i^{th} system call, a program calls `strace 1<<i`, where i is the syscall number (look in `kernel/syscall.h`).

You have to modify the xv6 kernel to print out a line when each system call is about to return if the system call's number is set in the mask.

The line should contain:

1. The process id
2. The name of the system call
3. The decimal value of the arguments(xv6 passes arguments via registers).
NOTE: You must always interpret the register value as an integer. (see kernel/syscall.c)
4. The return value of the syscall.

NOTE: The trace system call should enable tracing for the process that calls it and any children that it subsequently forks, but should not affect other processes.

Example:

```
$ strace 32 grep hello README
3: syscall read (1023 2736 1023) -> 1023
3: syscall read (968 2791 968) -> 968
3: syscall read (235 2741 1018) -> 235
3: syscall read (0 2736 1023) -> 0

$ strace 2147483647 grep hello README
3: syscall trace (0) -> 0
3: syscall exec (3 12208) -> 3
3: syscall open (3 0) -> 3
3: syscall read (1023 2736 1023) -> 1023
3: syscall read (968 2791 968) -> 968
3: syscall read (235 2741 1018) -> 235
3: syscall read (0 2736 1023) -> 0
3: syscall close (0) -> 0
```

HINTS:

- Add `$U/_strace` to `UPROGS` in `Makefile`
- Add a `sys_trace()` function in `kernel/sysproc.c` that implements the new system call by remembering its argument in a new variable in the `proc` structure (see `kernel/proc.h`). The functions to retrieve system call arguments from user space are in `kernel/syscall.c`, and you can see examples of their use in `kernel/sysproc.c`.

- Modify `fork()` (see `kernel/proc.c`) to copy the trace mask from the parent to the child process.
- Modify the `syscall()` function in `kernel/syscall.c` to print the trace output. You will need to add an array of syscall names to index into.
- Create a user program in `user/strace.c`, to generate the user-space stubs for the system call, add a prototype for the system call to `user/user.h`, a stub to `user/usys.pl`, and a syscall number to `kernel/syscall.h`. The Makefile invokes the Perl script `user/usys.pl`, which produces `user/usys.S`, the actual system call stubs, which use the RISC-V `ecall` instruction to transition to the kernel.

Specification 2: Scheduling [65 Marks]

The default scheduler of xv6 is round-robin-based. In this task, you'll implement 3 other scheduling policies and incorporate them in xv6. The kernel shall only use one scheduling policy which will be declared at compile time.

Modify the Makefile to support `SCHEDULER` - a macro for the compilation of the specified scheduling algorithm. Use the flags for compilation:-

- First Come First Serve = FCFS
- Priority Based = PBS
- Multilevel Feedback Queue = MLFQ

Example:

```
$ make qemu SCHEDULER=MLFQ
```

(a) First come - First Served (FCFS) [10 Marks]

Implement a policy that selects the process with the lowest creation time (creation time refers to the tick number when the process was created). The process will run until it no longer needs CPU time.

HINTS:

- a) Edit the `struct proc` (used for storing per-process information) in `kernel/proc.h` to store extra information about the process.

- b) Modify the `allocproc()` function to set up values when the process starts.
(see `kernel/proc.h`)
- c) Use preprocessor directives to declare the alternate scheduling policy in `scheduler()` in `kernel/proc.h`.
- d) Disable the preemption of the process after the clock interrupts in `kernel/trap.c`

(b) Priority Based Scheduler (PBS) [25 Marks]

Implement a non-preemptive priority-based scheduler that selects the process with the highest priority for execution. In case two or more processes have the same priority, we use the number of times the process has been scheduled to break the tie. If the tie remains, use the start-time of the process to break the tie (processes with lower start times should be scheduled further).

There are two types of priorities.

- The **Static Priority** of a process (SP) can be in the range [0,100], the smaller value will represent higher priority. Set the default priority of a process as 60. The lower the value the higher the priority.
- **Dynamic Priority** (DP) is calculated from static priority and niceness.
- The **niceness** is an integer in the range [0, 10] that measures what percentage of the time the process was sleeping (see `sleep()` in `kernel/proc.c`. xv6 allows a process to give up the CPU and sleep waiting for an event, and allows another process to wake the first process up)
- The meaning of **niceness** values are:
 - 5 is neutral
 - 10 helps priority by 5
 - 0 hurts priority by 5
- To calculate the niceness:
 - Record for how many ticks the process was sleeping and running from the last time it was scheduled by the kernel. (see `sleep()` & `wakeup()` in `kernel/proc.c`)
 - New processes start with niceness equal to 5. After scheduling the process, compute the niceness as follows:

$$\text{niceness} = \text{Int}\left(\frac{\text{Ticks spent in (sleeping) state}}{\text{Tick spent in (running + sleeping) state}} * 10 \right)$$

- Use Dynamic Priority to schedule processes which is given as:

$$DP = \max(0, \min(SP - \text{niceness} + 5, 100))$$

To change the Static Priority add a new system call **set_priority()**. This resets the niceness to 5 as well.

```
int set_priority(int new_priority, int pid)
```

The system call returns the old Static Priority of the process. **In case the priority of the process increases (the value is lower than before)**, then rescheduling should be done.

Also make sure to implement a user program **setpriority**, which uses the above system call to change the priority. And takes the syscall arguments as command-line arguments.

```
setpriority priority pid
```

NOTES:

- To implement `set_priority()` system call refer to the hints of specification 1 (Trace system call)
- Don't forget to update the niceness of the process.

(c) Multilevel Feedback queue scheduling (MLFQ) [30 Marks]

Implement a simplified preemptive MLFQ scheduler that allows processes to move between different priority queues based on their behavior and CPU bursts.

- If a process uses too much CPU time, it is pushed to a lower priority queue, leaving I/O bound and interactive processes in the higher priority queues.
- To prevent starvation, implement aging.

Details:

1. Create five priority queues, giving the highest priority to queue number 0 and lowest priority to queue number 4
2. The time-slice are as follows:
 - a. For priority 0: 1 timer tick
 - b. For priority 1: 2 timer ticks
 - c. For priority 2: 4 timer ticks
 - d. For priority 3: 8 timer ticks
 - e. For priority 4: 16 timer ticks

NOTE: Here tick refers to the clock interrupt timer. (see kernel/trap.c)

Synopsis for the scheduler:-

1. On the initiation of a process, push it to the end of the highest priority queue.
2. You should always run the processes that are in the highest priority queue that is not empty.

Example:

Initial Condition: A process is running in queue number 2 and there are no processes in both queues 1 and 0.

Now if another process enters in queue 0, then the current running process (residing in queue number 2) must be preempted and the process in queue 0 should be allocated the CPU.

3. When the process completes, it leaves the system.
4. If the process uses the complete time slice assigned for its current priority queue, it is preempted and inserted at the end of the next lower level queue.
5. If a process voluntarily relinquishes control of the CPU(eg. For doing I/O), it leaves the queuing network, and when the process becomes ready again after the I/O, it is inserted at the tail of the same queue, from which it is relinquished earlier (**Q: Explain in the README how could this be exploited by a process(see specification 4 below)**).
6. A round-robin scheduler should be used for processes at the lowest priority queue.
7. To prevent starvation, implement aging of the processes:
 - a. If the wait time of a process in a priority queue (other than priority 0) exceeds a given limit (assign a suitable limit to prevent starvation), their priority is increased and they are pushed to the next higher priority queue.
 - b. The wait time is reset to 0 whenever a process gets selected by the scheduler or if a change in the queue takes place (because of aging).

Specification 3: procdump [10 Marks]

procdump is a function that is useful for debugging (see `kernel/proc.c`). It prints a list of processes to the console when a user types Ctrl-p on the console. In this task, you have to extend this function to print more information about all the active processes. Check the sample output given below to know what all information needs to be displayed.

For MLFQ

PID	Priority	State	rtime	wtime	nrun	q0	q1	q2	q3	q4
1	2	sleeping	12	10	2	5	7	4	0	0
2	0	running	5	0	1	5	0	0	0	0
3	1	running	8	5	2	5	3	0	0	0
4	-1	zombie	7	8	3	1	2	4	8	0

For PBS

PID	Priority	State	rtime	wtime	nrun
1	60	sleeping	12	10	2
2	23	running	5	0	1
3	21	running	8	5	2
4	19	zombie	7	8	3

- **priority [PBS, MLFQ only]:**
 - PBS: Current dynamic-priority of the process. It will range from [0, 100].
 - MLFQ: This corresponds to the queue number of the process. Put -1 in case the process is currently not queued in any of the queues.
- **state:** The current state of the process. (see enum `procstate` in `kernel/proc.h`)

- **rtime:** Total ticks for which the process ran on the CPU till now.
- **wtime:** This corresponds to the total waiting time for the process. However, In the case of the MLFQ scheduler, this is the wait time in the current queue.
- **nrun:** Number of times the process was picked by the scheduler.
- **q_i[MLFQ only]:** Number of ticks the process has spent in each of the 5 queues.

Specification 4: Report [10 Marks]

- Include well-written documentation of how you implemented the other 3 specifications.
- Answer the question asked in specification 2, subtopic MLFQ with proper reasoning.
- We'll provide a sample benchmark program. Tabulate the performances of the scheduling algorithms using the given benchmark program (or implement your own). A similar benchmark program will be used to demonstrate the same during the evals.
- Include the performance comparison between the default and 3 implemented policies in the README by showing the average waiting and running times for processes. (You may find the use of `waitx()` syscall implemented in the tutorial to be quite helpful)


Bonus (Optional): [10 Marks]

MLFQ scheduling analysis

Plot timeline graphs for processes running with MLFQ Scheduler. Use the benchmark from Specification 2 to vary how long each process uses the CPU before relinquishing voluntarily (Hint: use `sleep()`). The graph should be a timeline/scatter plot between `queue_id`(y-axis) and `time elapsed`(x-axis) from start with color-coded processes.

Add to the README the observations recorded for different types of processes.

Example:



Guidelines

1. Submission format: <RollNo>_Assignment4.tar.gz.
2. Submission by email to TAs will not be accepted.
3. Any copy cases found will lead to serious consequences.
4. **We will use more than 2 CPUs to test for FCFS and PBS. But for the MLFQ scheduler, we will only use 1 CPU.**
5. Make sure you write a README that contains the report components of the assignment. Including a README file is *NECESSARY*
6. Whenever you add new files do not forget to add them to the Makefile so that they get included in the build.